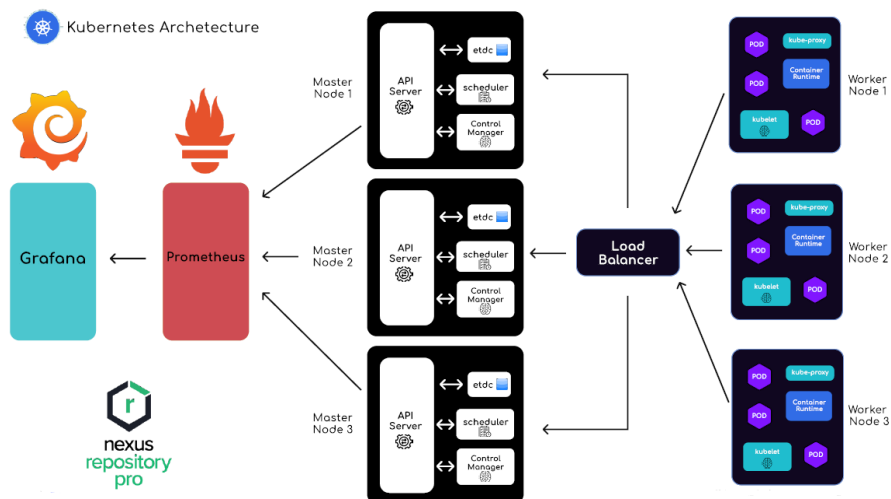


[5조] 내부 etcd Workbook

- Ubuntu 20.04 LTS
- Docker 20.10.18
- Kubernetes v1.24.5
- memory : 4GB
- disk : 100GB
- /dev/sda1 xfs / : 70GB
- /dev/sda5 swap : 8GB
- /dev/sda6 xfs /var/lib/docker : 29372MB (나머지)

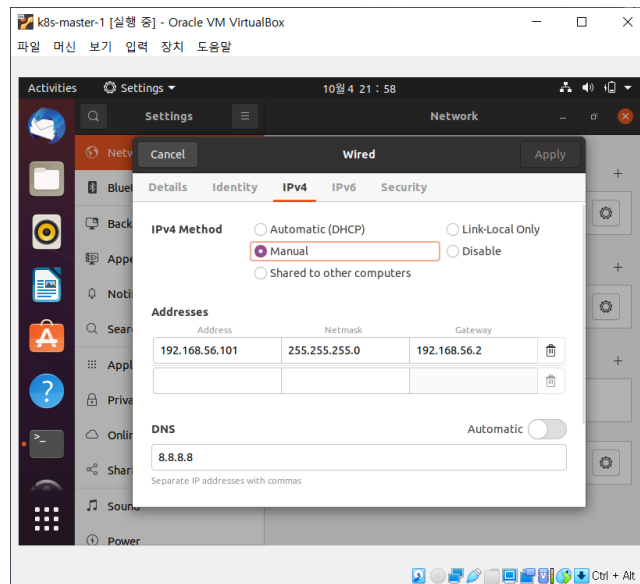
Stacked etcd HA Cluster Topology



1 서버 셋팅

IP 설정

enp0s8 설정 → 설정 후 enp0s8 한번 꺾다 쳐주기



📌 기본 설치 및 설정

```
# hostname 설정
kevin@master-1:~$ sudo hostnamectl set-hostname master-1

# vim, openssh, net-tools 설치
kevin@master-1:~$ sudo apt -y update
kevin@master-1:~$ sudo apt -y install vim
kevin@master-1:~$ sudo apt -y install openssh-server
kevin@master-1:~$ sudo apt install net-tools
kevin@master-1:~$ sudo reboot

##### putty 연결 #####

# 방화벽 중지
kevin@master-1:~$ sudo apt -y install firewalld
kevin@master-1:~$ sudo systemctl daemon-reload
kevin@master-1:~$ sudo systemctl stop firewalld.service
kevin@master-1:~$ sudo firewall-cmd --reload
FirewallD is not running

# swap off
kevin@master-1:~$ sudo vi /etc/fstab
kevin@master-1:~$ sudo -i
root@master-1:~# swapoff -a
root@master-1:~# echo 0 > /proc/sys/vm/swappiness
root@master-1:~# sed -e '/swap/ s/^#*#/' -i /etc/fstab
root@master-1:~# exit
logout

# 아래와 같이 swap 주석처리 되어있는지 확인
kevin@master-1:~$ sudo vi /etc/fstab
# swap was on /dev/sda5 during installation
#UUID=7d3b72c7-2162-4d28-9210-dae2233515f9 none swap sw 0 0

# ntp(network time protocole) - 시간동기화
kevin@master-1:~$ sudo apt -y install ntp
kevin@master-1:~$ sudo systemctl daemon-reload
kevin@master-1:~$ sudo systemctl enable ntp
kevin@master-1:~$ sudo systemctl restart ntp
kevin@master-1:~$ sudo ntpq -p
kevin@master-1:~$ date

# network forward 설정
# why? docker private registry를 이용하기 위해서 적용
kevin@master-1:~$ sudo su -
root@master-1:~# echo '1' > /proc/sys/net/ipv4/ip_forward
root@master-1:~# cat /proc/sys/net/ipv4/ip_forward
1 # 1로 나와야함 !! 0이면 안됨
root@master-1:~# exit
logout
```



스왑 메모리를 비활성화 해야하는 이유

- Pod를 할당하고 제어하는 kubelet은 스왑 상황을 처리하도록 설계되지 않음 !
- why? kubernetes에서 가장 기본이 되는 Pod의 컨셉이 필요한 리소스 만큼만 호스트 자원에서 할당 받아 사용한다.
- 따라서, kubernetes 개발팀이 메모리 스왑을 고려하지 않고 설계했기 때문에 클러스터 노드로 사용할 서버 머신들은 모두 스왑 메모리를 비활성화 해줘야 한다.



부팅 시 시간 동기화 해결 ::: 부팅 시 자동으로 rdate 명령어 실행

vi /etc/rc.d/rc.local 파일 아래에 rdate -s time.bora.net 명령을 적어준다.

이 과정을 통해 매번 부팅 시마다 자동으로 시간을 동기화할 수 있다.

😊 만약 rc.local이 존재하지 않는다면(ubuntu 18버전 이후) 아래 블로그를 참조하세요

[리눅스 유용한 팁] 우분투(Ubuntu)에서 rc.local 없을 때, 활성화 시키는 방법!

이번 포스팅에서는 리눅스를 사용하면서 새롭게 부팅될 때 마다 자동으로 먼가를 실행하도록 하고 싶을 때,...

<https://m.blog.naver.com/PostView.naver?isHttpsRedirect=true&blogId=sw4r&logNo=221744290073>

```
root@porcupine-pc:~# systemctl status rc-local.service
rc-local.service - /etc/rc.local Compatibility
Loaded: loaded (/lib/systemd/system/rc-local.service; enabled; vendor preset: enabled)
Drop-In: /lib/systemd/system/rc-local.service.d
~debian.conf
Active: active (exited) since Fri 2019-12-20 13:45:47 GMT; 1h 52min ago
Docs: man:systemd-rc-local-generator(8)
Tasks: 0 (limit: 4031)
Memory: 0K
CGroup: /system.slice/rc-local.service
Dec 20 13:45:47 porcupine-pc systemd[1]: Starting /etc/rc.local Compatibility...
Dec 20 13:45:47 porcupine-pc systemd[1]: Started /etc/rc.local Compatibility.
```

2 로드밸런서 서버 설치

```
kevin@master-1:~$ sudo apt update
kevin@master-1:~$ sudo apt install -y apache2
kevin@master-1:~$ sudo apt install -y haproxy
kevin@master-1:~$ sudo vi /etc/haproxy/haproxy.cfg
global
...

defaults
log      global
mode     tcp      # https 트래픽도 받을 수 있도록 tcp로 설정
option   tcplog    # https 트래픽도 받을 수 있도록 tcp로 설정
...

frontend proxynode
bind *:6443
stats uri /proxystats
default_backend masters

backend masters
balance roundrobin
server master-1 192.168.56.101:6443 check # master1
server master-2 192.168.56.102:6443 check # master2
server master-3 192.168.56.103:6443 check # master3

listen stats
bind :9999
mode http
stats enable
stats hide-version
stats uri /stats

kevin@master-1:~$ sudo systemctl stop apache2

# haproxy를 restart 하여 haproxy.cfg 설정내용을 적용
kevin@master-1:~$ sudo systemctl restart haproxy
kevin@master-1:~$ sudo systemctl status haproxy
active (running)

kevin@master-1:~$ netstat -nltp | grep 6443
(Not all processes could be identified, non-owned process info
will not be shown, you would have to be root to see it all.)
tcp        0      0 0.0.0.0:0:6443          0.0.0.0:*               LISTEN      -
```

로드밸런서 확인 사항

```

kevin@lb:~$ netstat -nltp | grep 6443
(Not all processes could be identified, non-owned process info
 will not be shown, you would have to be root to see it all.)
tcp        0      0 0.0.0.0:6443        0.0.0.0:*           LISTEN      -

kevin@lb:~$ nc -v lb 6443
Connection to lb 6443 port [tcp/*] succeeded!

```

HProxy에서 마스터 노드 접속 상태 확인하기

로드밸런서의 9999번 포트로 접속하면 HAproxy 동작 상태를 확인할 수 있다.

<http://192.168.56.100:9999/stats>

Statistics Report for HAPROXY

192.168.56.100:9999/stats

네이버 MYBOX

HAProxy

Statistics Report for pid 2575

> General process information

pid = 2575 (process #1, nbproc = 1, nthread = 4)
uptime = 0d 0h22m49s
system limits: memmax = unlimited; ulimit-n = 524288
maxsock = 524288; **maxconn** = 262120; **maxpipes** = 0
current conn = 6; **current** pipes = 0/0; **conn rate** = 1/sec; **bit rate** = 0.815 kbps
Running tasks: 1/2s; **idle** = 100 %

active UP
active UP, going down
active DOWN, going up
active or backup DOWN
active or backup DOWN for maintenance (MAINT)
active or backup SOFT STOPPED for maintenance
Note: "NOLE"/"DRAIN" = UP with load-balancing disabled.

backup UP
backup UP, going down
backup DOWN, going up
not checked

Display option:

- Scope:
- [Hide "DOWN" servers](#)
- [Refresh now](#)
- [CSV export](#)

External resources:

- [Primary site](#)
- [Updates \(v2.0\)](#)
- [Online manual](#)

proxynode		Queue			Session rate			Sessions					Bytes				Denied		Errors			Warnings		Server										
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwntime	Thrtle			
Frontend		0	0	-	0	11	-	4	16	262	120	194			238	403	1	027	283	0	0	0			OPEN									

masters		Queue			Session rate			Sessions					Bytes				Denied		Errors			Warnings		Server											
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwntime	Thrtle				
master-1		0	0	-	0	8	-	4	16	-	53	48	31s	235	264	1	027	283	0	0		0	5	0	4s	DOWN 1/2	L4OK in 0ms	1	Y	-	188	22	4m30s	-	
Backend		0	0		0	11		4	16	26	212	194	48	31s	238	403	1	027	283	0	0		146	5	0	4s	DOWN		0	0	0		22	4m30s	

stats		Queue			Session rate			Sessions					Bytes				Denied		Errors			Warnings		Server											
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Dwntime	Thrtle				
Frontend		1	5	-	2	5	-	2	5	262	120	50		32	032	619	259	0	0	4						OPEN									
Backend		0	0		1	4		0	1	26	212	43	0	0s	32	032	619	259	0	0		43	0	0	0	22m49s	UP		0	0	0		0		

master-1 kubaem init 후

Statistics Report for HAProxy

주요 요약 | 192.168.56.100:9999/stats

네이버 MYBOX

HAProxy

Statistics Report for pid 3449

> General process information

pid = 3449 (process #1, nproc = 1, nbthread = 4)
 uptime = 0d 0h00m08s
 system limits: memmax = unlimited, ulimit-n = 524287
 maxsock = 524287, maxconn = 262119, maxpipes = 0
 current conns = 14, current pipes = 0/0, conn rate = 2/sec, bit rate = 0.000 kbps
 Running tasks: 1/38, idle = 100 %

active UP, active UP, going down, active DOWN, going up, active or backup DOWN, active or backup DOWN for maintenance (MAINT), active or backup SOFT STOPPED for maintenance, backup UP, backup UP, going down, backup DOWN, going up, not checked

Display option: Scope:
 Hide 'DOWN' servers, Refresh now, CSV export

External resources: Primary site, Updates (v2.0), Online manual

Note: "NOLE"/"DRAIN" = UP with load-balancing disabled.

proxynode		Queue			Session rate			Sessions				Bytes		Denied		Errors		Warnings		Server												
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Downtime	Thrtle	
Frontend		0	18	-	12	18		262	119	18				6	272	9	588	0	0	0	0	0	0	OPEN								

masters		Queue			Session rate			Sessions				Bytes		Denied		Errors		Warnings		Server											
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Downtime	Thrtle
master-1		0	0	-	0	9		3	9		9	9	7s	6	272	9	588	0	0	0	0	0	8s UP	L4OK in 0ms	1	Y	-	0	0	0s	-
master-2		0	0	-	0	9		8	9	-	19	9	2s	0	0	0	0	0	0	11	0	5s DOWN	L4CON in 1024ms	1	Y	-	1	1	5s	-	
Backend		0	0		0	18		12	18		26	212	18	18	7s	6	272	9	588	0	0	0	8s UP		1	1	0		0	0s	

stats		Queue			Session rate			Sessions				Bytes		Denied		Errors		Warnings		Server											
		Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	Last	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Downtime	Thrtle
Frontend		2	2	-	2	2		262	119	2				0	0	0	0	0	0	0	0	0	OPEN								
Backend		0	0		0	0		0	0		26	212	0	0	0s	0	0	0	0	0	0	0	8s UP		0	0	0		0		

master-2 join 후

3 master node && worker node 공통

도커 설치

```

kevin@master-1:~$ sudo apt-get install -y apt-transport-https ca-certificates curl software-properties-common gnupg2

# curl 명령어를 통해 gpg key 내려받기
kevin@master-1:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key --keyring /etc/apt/trusted.gpg.d/docker.gpg add
OK

# 패키지 관리 도구에 도커 다운로드 링크 추가
kevin@master-1:~$ sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"

# focal 확인
kevin@master-1:~$ tail /etc/apt/sources.list
deb [arch=amd64] https://download.docker.com/linux/ubuntu focal stable

# 패키지 관리 도구 업데이트
kevin@master-1:~$ sudo apt -y update

# 도커 설치
kevin@master-1:~$ sudo apt-get -y install docker-ce

# 설치 확인
kevin@master-1:~$ sudo docker version
kevin@master-1:~$ sudo docker info | grep cgroup
Cgroup Driver: cgroupfs

# Docker 데몬이 사용하는 드라이버를 cgroupfs 대신 systemd를 사용하도록 설정
# why? kubernetes에서 권장하는 Docker 데몬의 드라이버는 systemd 이고,
# systemd를 사용하면 kubernetes가 클러스터 노드에서 사용 가능한 자원을 쉽게 알 수 있도록 해준다.
kevin@master-1:~$ sudo vi /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2",
}

```

```

"storage-opts": [
  "overlay2.override_kernel_check=true"
]
}
kevin@master-1:~$ sudo mkdir -p /etc/systemd/system/docker.service.d
kevin@master-1:~$ sudo systemctl daemon-reload
kevin@master-1:~$ sudo systemctl enable docker
kevin@master-1:~$ sudo systemctl restart docker

kevin@master-1:~$ sudo systemctl status docker
active (running)
kevin@master-1:~$ sudo docker info | grep -i cgroup
Cgroup Driver: systemd
Cgroup Version: 1

```



config.toml 파일명 변경

containerd를 CRI 런타임으로 사용하기 위함

```

kevin@master-1:~$ cd /etc/containerd/
kevin@master-1:/etc/containerd$ sudo mv config.toml config.toml.org
kevin@master-1:/etc/containerd$ ls
config.toml.org

kevin@master-1:/etc/containerd$ cd
kevin@master-1:~$ sudo systemctl restart containerd.service
kevin@master-1:~$ sudo systemctl restart kubelet

```



쿠버네티스 설치

```

kevin@master-1:~$ curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-key add -
OK
kevin@master-1:~$ cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
> deb https://apt.kubernetes.io/ kubernetes-xenial main
> EOF
deb https://apt.kubernetes.io/ kubernetes-xenial main
kevin@master-1:~$ sudo apt -y update

kevin@master-1:~$ sudo apt -y install kubeadm=1.24.5-00 kubelet=1.24.5-00 kubectl=1.24.5-00
kevin@master-1:~$ sudo apt list | grep kubernetes
...
kubeadm/kubernetes-xenial 1.25.2-00 amd64 [upgradable from: 1.24.5-00]
kubectl/kubernetes-xenial 1.25.2-00 amd64 [upgradable from: 1.24.5-00]
kubelet/kubernetes-xenial 1.25.2-00 amd64 [upgradable from: 1.24.5-00]
...

kevin@master-1:~$ sudo apt-mark hold kubelet kubeadm kubectl
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.

```



hosts 파일 설정

```

kevin@master-1:~$ sudo vi /etc/hosts
127.0.0.1        localhost
127.0.1.1        master-1          # 복제 후 각각 변경해야 한다.
192.168.56.100   lb
192.168.56.101   master-1
192.168.56.102   master-2
192.168.56.103   master-3
192.168.56.201   worker-1
192.168.56.202   worker-2
192.168.56.203   worker-3

```



복제 후 확인 사항

- 모든 네트워크 어댑터의 새 MAC 주소 생성 , 완전한 복제 로 복제하기
- IP 및 hostname 변경
- putty 연결
- sudo vi /etc/hosts 수정
- 노드 간 ping 확인
- 방화벽 확인

```
kevin@master-3:~$ sudo systemctl daemon-reload
kevin@master-3:~$ sudo systemctl stop firewalld.service
kevin@master-3:~$ sudo firewall-cmd --reload
Firewalld is not running
```

- date 확인
- 모든 노드끼리 ssh 접근
 - 노드 간 원격접속 되는지 확인

```
kevin@master-1:~$ ssh k8s-node1
kevin@master-1:~$ ssh k8s-node2
...

# 확인
kevin@master-1:~$ cat .ssh/known_hosts
|1|c7ewwIBvjFvAN9lw9c5QiH8H5N8=|g3NBTvsjFvzLQ1TLTRX6ptAsw3Y= ecdsa-sha2-nistp256 AAAAE2VjZHNhLXNoYTItbmlzdHAyNTYAAAAIbmlzdHAyNTYAAQ==
|1|WnwoezAVRfLwFh+sYfVZVBMp2DU=|xH7GN6SIDZRWE1VimN7uzJPkd4= ecdsa-sha2-nistp256 AAAAE2VjZHNhLXNoYTItbmlzdHAyNTYAAAAIbmlzdHAyNTYAAQ==
|1|EUFNLCL5outdMqSqC4qs2AgaYvQ=|vYu2q4Z1uiP75cQy0ZcDhF1r57Q= ecdsa-sha2-nistp256 AAAAE2VjZHNhLXNoYTItbmlzdHAyNTYAAAAIbmlzdHAyNTYAAQ==
|1|ju2Gw1WHYNfcNa7WgSNaXflrowY=|4SaaorFJfjoQkygHoKJli0RCu6Y= ecdsa-sha2-nistp256 AAAAE2VjZHNhLXNoYTItbmlzdHAyNTYAAAAIbmlzdHAyNTYAAQ==
```

- 쿠버네티스 버전 확인

```
kevin@master-1:~$ kubeadm version
kubeadm version: &version.Info{Major:"1", Minor:"24", GitVersion:"v1.24.5", GitCommit:"e979822c185a14537054f15808a118d7fcce1d6e",
kevin@master-1:~$ kubelet --version
Kubernetes v1.24.5
```

4 kubernetes 초기화 작업

(master-1 에서) kubeadm init

```
kevin@master-1:~$ sudo kubeadm init --control-plane-endpoint=lb:6443 --apiserver-advertise-address=192.168.56.101 --pod-network-cidr=10.244.0.0/16
...
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

You can now join any number of the control-plane node running the following command on each as root:

kubeadm join lb:6443 --token ap62c8.tz1v3xmu1v20pju3 \
--discovery-token-ca-cert-hash sha256:1d611f9a4ee26405f81cfd578234708a28e6b8c33924a8bad2136131f7ae49b \
--control-plane --certificate-key 9628886c1cd06f5d1819a14771c36086fd63e660a90baa193e622fa8bfb4c892

Please note that the certificate-key gives access to cluster sensitive data, keep it secret!
As a safeguard, uploaded-certs will be deleted in two hours; If necessary, you can use
"kubeadm init phase upload-certs --upload-certs" to reload certs afterward.
```

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join lb:6443 --token ap62c8.tz1v3xmu1v20pju3 \
--discovery-token-ca-cert-hash sha256:1d611f9a4ee26405f81cfd578234708a28e6b8c33924a8bad2136131f7ae49b
```

kubeadm init 옵션

- `--control-plane-endpoint`
 - 모든 control-plane에 대한 공유 엔드 포인트를 설정하는 옵션이다. control-plane이 바라보는 주소를 적는다. 이 옵션을 통해 로드밸런서 기능을 사용할 수 있다. 마스터 노드 클러스터 생성 시 필요하다.
 - 공식 문서에서는 해당 옵션에 ip address를 그대로 사용하지 말고, DNS name을 이용하길 권장한다.

Creating Highly Available Clusters with kubeadm

This page explains two different approaches to setting up a highly available Kubernetes cluster using kubeadm: With stacked control plane nodes. This approach requires less infrastructure. The etcd members and control plane nodes are co-located. With an external etcd cluster. This approach requires

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/high-availability/>

kubern

- `--apiserver-advertise-address`
 - control-plane API Server의 IP를 설정한다.
- `--pod-network-cidr`
 - pod의 네트워크를 설정할 때 사용
- `--service-cidr`
 - pod와 service의 네트워크 대역 분리를 위해 사용
- `--upload-certs`
 - MasterNode1의 Certificate가 다른 노드에 자동 배포되도록 한다. 마스터 노드 클러스터 생성 시 필요하다.

```
kevin@master-1:~$ mkdir -p $HOME/.kube
kevin@master-1:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
kevin@master-1:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

확인

```
kevin@master-1:~$ kubectl get no
NAME      STATUS    ROLES    AGE    VERSION
master-1  NotReady  master   10m    v1.19.15
```

📌 (master-2, master-3 에서) control plane join

```
kevin@master-2:~$ sudo kubeadm join lb:6443 --token ap62c8.tz1v3xmu1v20pju3 \
> --discovery-token-ca-cert-hash sha256:1d611f9a4ee26405f81cfd578234708a28e6b8c33924a8bad2136131f7ae49b \
> --control-plane --certificate-key 9628886c1cd06f5d1819a14771c36086fd63e660a90baa193e622fa8bfb4c892
```

...

This node has joined the cluster and a new control plane instance was created:

- * Certificate signing request was sent to apiserver and approval was received.
- * The Kubelet was informed of the new secure connection details.
- * Control plane (master) label and taint were applied to the new node.
- * The Kubernetes control plane instances scaled up.
- * A new etcd member was added to the local/stacked etcd cluster.

To start administering your cluster from this node, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Run 'kubectl get nodes' to see this node join the cluster.


```
# kubectl 명령어 사용을 위해 아래 명령어를 실행
kevin@master-2:~$ mkdir -p $HOME/.kube
kevin@master-2:~$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
kevin@master-2:~$ sudo chown $(id -u):$(id -g) $HOME/.kube/config

# kubectl get nodes로 클러스터가 잘 생성되었는지 확인
```

worker node join

```
kevin@worker-2:~$ kubeadm join lb:6443 --token ap62c8.tz1v3xmu1v20pju3 \
--discovery-token-ca-cert-hash sha256:1d611f9a4ee26405f81cfda578234708a28e6b8c33924a8bad2136131f7ae49b
```

5 k8s에서 서비스 올리기

- `react` : `v16`
- `django` : `v4.1`
- `MySQL` : `v8.0`



아래에 해당하는 포트는 제외하고 서비스를 올려합니다.

- 마스터 노드에서 사용하는 포트
 - 6443 포트 : Kubernetes API Server / Used By All
 - 2379~2380 포트 : etcd server client API / Used By kube-apiserver, etcd
 - 10250 포트 : Kubelet API / Used By Self, Control plane
 - 10251 포트 : kube-scheduler / Used By Self
 - 10252 포트 : kube-controller-manager / Used By Self
- 워커 노드에서 사용하는 포트
 - 10250 포트 : Kubelet API / Used By Self, Control plane
 - 30000~32767 포트 : NodePort Services / Used By All

Dockerfile 생성 후 이미지 빌드하기

DB

```
FROM mysql:8.0

WORKDIR /

COPY db.sql /docker-entrypoint-initdb.d

ENV MYSQL_ROOT_PASSWORD=password
ENV MYSQL_DATABASE=Nexpot

EXPOSE 3306

CMD ["mysqld"]
```

Backend

```
FROM python:3.9

WORKDIR /

COPY requirements.txt ./
RUN /usr/local/bin/python -m pip install --upgrade pip
RUN pip install --no-cache-dir -r requirements.txt

COPY . /project
WORKDIR /project

EXPOSE 8000

CMD [ "python", "manage.py", "runserver", "0.0.0.0:8000" ]
```

Frontend

```
FROM node:18

WORKDIR /app
COPY ./front-end .
RUN npm install

EXPOSE 3000

CMD [ "npm", "start", "--host", "0.0.0.0" ]
```

k8s를 위한 Manifest 파일 작성

db-pv-pvc.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: db-pv
  labels:
    app: db-pv
spec:
  storageClassName: manual
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: "/mnt/mysql" ## 여러 가지 기타의 파일 시스템을 마운트 하는 곳
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: db-pvc
spec:
  storageClassName: manual
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

db-deploy-svc.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: db-deploy
  labels:
    app: db
spec:
  replicas: 1
  selector:
    matchLabels:
```

```

    app: db
  template:
    metadata:
      labels:
        app: db
    spec:
      containers:
        - name: db-container
          image: zxcasd3004/app:db1
          ports:
            - containerPort: 3306
          volumeMounts:
            - name: db-vol
              mountPath: /var/lib/mysql
      volumes:
        - name: db-vol
          persistentVolumeClaim:
            claimName: db-pvc
---
apiVersion: v1
kind: Service
metadata:
  name: db-svc
spec:
  selector:
    app: db
  ports:
    - port: 3333
      targetPort: 3306

```

backend-deploy-svc.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: be-deploy
  labels:
    app: be
spec:
  replicas: 3
  selector:
    matchLabels:
      app: be
  template:
    metadata:
      labels:
        app: be
    spec:
      containers:
        - name: be-container
          image: ur2e/be:1.0
          ports:
            - containerPort: 8000
          env:
            - name: DATABASE
              value: db-svc
---
apiVersion: v1 # frontend (react)에게 request보낼 backend 파드의 ip 주소를 알려주기 위해 be-svc 생성
kind: Service
metadata:
  name: be-svc
spec:
  selector:
    app: be
  ports:
    - port: 8888
      targetPort: 8000

```

fe-secret.yaml

```

# kakao API
apiVersion: v1
kind: Secret
metadata:
  name: fe-secret
data:
  REACT_APP_KAKAOMAP_API_KEY: d413cb240#$$@#^@#^@#6166

```

frontend-deploy-svc.yaml

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: fe-deploy
  labels:
    app: fe
spec:
  replicas: 1
  selector:
    matchLabels:
      app: fe
  template:
    metadata:
      labels:
        app: fe
    spec:
      nodeSelector:
        kubernetes.io/hostname: k8s-node2
      containers:
        - name: fe-container
          image: ur2e/fe:4.0
          ports:
            - containerPort: 3000
          env: ## BACKEND와 통신하기 위한 환경변수 설정
            - name: BACKEND
              value: be-svc
          volumeMounts: ## react에게 Kakao API Key를 알려주기위함.
            - name: api-key
              mountPath: "/app/.env"
              readOnly: true
      volumes:
        - name: api-key
          secret:
            secretName: fe-secret
---
apiVersion: v1
kind: Service
metadata:
  name: fe-svc
spec:
  selector:
    app: fe
  ports:
    - port: 3000
      targetPort: 3000
  externalIPs:
    - 192.168.56.102

```

모든 yaml 파일 apply 후 pod와 service 조회

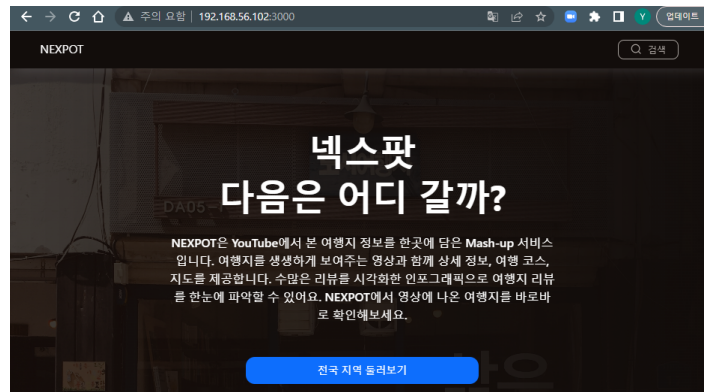
\$ kubectl get po,svc,deploy

NAME	READY	STATUS	RESTARTS	AGE
pod/be-deploy-6484676968-28kmb	1/1	Running	0	10h
pod/be-deploy-6484676968-kxt2t	1/1	Running	0	10h
pod/be-deploy-6484676968-lg7b5	1/1	Running	0	10h
pod/db-deploy-d6dcb9bdb-vjs68	1/1	Running	0	10h
pod/fe-deploy-5b9b99c65-cd8ff	1/1	Running	0	10h
pod/fe-deploy-5b9b99c65-jng6v	1/1	Running	0	10h
pod/fe-deploy-5b9b99c65-tq5c5	1/1	Running	0	10h
pod/nginx	1/1	Running	0	13h
pod/nginx-m2	1/1	Running	0	13h
pod/nginx-m3	1/1	Running	0	13h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/be-svc	ClusterIP	20.97.89.246	<none>	8888/TCP	10h
service/db-svc	ClusterIP	20.100.75.230	<none>	3333/TCP	10h
service/fe-svc	ClusterIP	20.98.255.204	192.168.56.202	3000/TCP	10h
service/kubernetes	ClusterIP	20.96.0.1	<none>	443/TCP	4d9h
service/nginx	NodePort	20.97.146.85	<none>	80:30273/TCP	13h
service/nodep-svc	LoadBalancer	20.96.148.127	<pending>	9000:30057/TCP	34h

NAME	READY	UP-T0-DATE	AVAILABLE	AGE
deployment.apps/be-deploy	3/3	3	3	10h
deployment.apps/db-deploy	1/1	1	1	10h
deployment.apps/fe-deploy	3/3	3	3	10h

📌 service/fe-svc의 External-ip:3000 으로 접속



6 prometheus + grafana 모니터링

마스터 노드 ⇒ 익스포터 ⇒ 프로메테우스 ⇒ 그라파나 ⇒ 브라우저

```
mkdir k8s-prometheus && cd $_
git clone https://github.com/brayanlee/k8s-prometheus.git

# Prometheus
kubectl create namespace monitoring
kubectl create -f prometheus/prometheus-ConfigMap.yaml
kubectl create -f prometheus/prometheus-ClusterRoleBinding.yaml
kubectl create -f prometheus/prometheus-ClusterRole.yaml
kubectl create -f prometheus/prometheus-Deployment.yaml
kubectl create -f prometheus/prometheus-Service.yaml
kubectl create -f prometheus/prometheus-DaemonSet-nodeexporter.yaml

# kube-state
kubectl create -f kube-state/kube-state-ClusterRoleBinding.yaml
kubectl create -f kube-state/kube-state-ClusterRole.yaml
kubectl create -f kube-state/kube-state-Service
kubectl create -f kube-state/kube-state-ServiceAccount.yaml
kubectl create -f kube-state/kube-state-Deployment.yaml
kubectl create -f kube-state/kube-state-Service.yaml

# grafana
kubectl create -f grafana/grafana-Deployment.yaml
kubectl create -f grafana/grafana-Service.yaml
```

+ etcd 데이터 공유

- version
 - `etcdctl` : 3.2.26
 - `API` : 3.2

```
# etcd 설치
sudo apt-get -y update && sudo apt-get -y install etcd-client

# 버전 확인
yji@k8s-master:~$ etcdctl version
etcdctl version: 3.2.26
API version: 3.2

/*
 * 📌 실습 1. etcd에 직접 데이터 생성 후 다른 노드에서 조회하기
 */

📌 master 1
# etcd3ctl alias 설정
alias etcd3ctl="ETCDCTL_API=3 etcdctl --endpoints=https://192.168.56.101:2379 --cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kub
```

```

root@k8s-master-1:~# etcdctl put key1 value1
OK
root@k8s-master-1:~# etcdctl get key1 value1
key1
value1

❏ master 2
# etcdctl alias 설정
alias etcd3ctl="ETCDCTL_API=3 etcdctl --endpoints=https://192.168.56.102:2379 --cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kub

root@k8s-master-1:~# etcdctl member list
181ffe2d394f445b, started, k8s-master-1, https://192.168.56.101:2380, https://192.168.56.101:2379
7dcbf96552ea4f10, started, k8s-master-3, https://192.168.56.103:2380, https://192.168.56.103:2379
995e4e0f6e68ae86, started, k8s-master-2, https://192.168.56.102:2380, https://192.168.56.102:2379

# master 2에서 master 1에서 넣은 데이터 조회
root@k8s-master-2:~# etcdctl put etcd-test1 I-am-master-1
OK
root@k8s-master-2:~# etcdctl get etcd-test1
etcd-test1
I-am-master-1

❏ master 3
alias etcd3ctl="ETCDCTL_API=3 etcdctl --endpoints=https://192.168.56.103:2379 --cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kub
# master 3에서 master 1에서 넣은 데이터 조회
root@k8s-master-3:~# etcdctl put etcd-test1 I-am-master-1
OK
root@k8s-master-3:~# etcdctl get etcd-test1
etcd-test1
I-am-master-1

-----
/*
* ■ 실습 2. pod 생성 후 etcd에 저장된 파드 확인하기
*/

❏ master 1
# nginx-m2라는 이름의 pod 생성
$ kubectl run nginx --image=nginx-m2

# etcd를 조회하기위해 관리자 모드로 변경
$ su -

# etcd에 저장된 데이터를 알아보기위해 etcdctl get 명령을 json 확장자 파일로 redirect
$ etcdctl get / --prefix=true -w json > etcd_data.json

# jq 패키지를 사용해서 정렬되지 않은 json 파일을 포매팅 한다.
$ jq '.' etcd_data.json

- jq 패키지 실행 전 etcd_data.json
{"header":{"cluster_id":"15516845868796936735","member_id":"11051356363797606022","revision":517184,"raft_term":12},"kvs":[{"key":"L3J
json...}}]}

- jq 패키지 실행 후 etcd_data.json // etcd는 key-value 형태의 스토리지를 확인할 수 있음
{
  "header": {
    "cluster_id": 15516845868796936000,
    "member_id": 11051356363797606000,
    "revision": 517184,
    "raft_term": 12
  },
  "kvs": [
    {
      "key": "L3JlZ2lzdHJ5L2FwaWV4dGVuc2lvbnMuazhzLmVlL2N1c3RvbXJlc291cmNlZGVmaw5pdGVbnMvYmVmdWY29uZmVndXJhdGVbnMuY3JkLnByb2p1Y3Rj
      "create_revision": 7146,
      "mod_revision": 7150,
      "version": 3,
      "value": "eyJraW5kIjojQ3VzdG9tUmVzb3VyY2VEZWZpbml0aw9uIiwiaXBiYmVyc2lvbiI6ImFwaWV4dGVuc2lvbnMuazhzLmVlL3YxYmV0YT"
    }
  ]
}

# key 값 디코딩하여 etcd에 저장된 데이터를 확인
# jq '.kvs[].key' : jq로 원하는 kvs키의 value 값
# cut -d ' ' -f2: "(쌍따옴표)를 구분자로 2째 필드값을 자르기
# base64 --decode: base64-decoding
$ for k in $(cat etcd_data.json | jq '.kvs[].key' | cut -d ' ' -f2); do echo $k | base64 --decode; echo; done > etcd-data.txt

# 디코딩된 값 조회
$ vi etcd-data.txt
/registry/csinodes/k8s-master-1
/registry/csinodes/k8s-master-2
/registry/csinodes/k8s-master-3
/registry/daemonsets/kube-system/calico-node
/registry/daemonsets/kube-system/kube-proxy

```

```
registry/daemonsets/node-exporter
/registry/pods/default/nginx-m1
/registry/pods/default/nginx-m2 # 아까 생성한 파드 nginx-m2가 저장

# 아까 생성한 pod를 etcdctl 명령으로 조회하여 etcd가 저장한 pod의 정보를 조회
# 문자열 데이터가 아닌 부분은 깨져있다.
$ etcdctl get /registry/pods/default/nginx-m2 > etcd-nginx-m2
^Bv1^R^Rmetadata.namespace^P^C^R<91>^A
^H^Rnginx-m2^R^Enginx*^@B^@JL
^UKube-api-access-hvvvg^PA^Z-/var/run/secrets/kubernetes.io/serviceaccount"^@2^@j^T/dev/termination-logr^FAlways<80>^A^@<88>^A^@<90>^A
^GRunning^R#
^KInitialized^R^DTrue^Z^@"^HHäã<93><9a>^F^P^@*^@2^@^R^]
^EReady^R^DTrue^Z^@"^HHäã<93><9a>^F^P^@*^@2^@^R'
^OContainersReady^R^DTrue^Z^@"^HHäã<93><9a>^F^P^@*^@2^@^R$
^LPodScheduled^R^DTrue^Z^@"^HH<81>><0><93><9a>^F^P^@*^@2^@^Z^@"^@*^N192.168.56.2032^L10.98.69.216:^HHäã<93><9a>^F^P^@Bä^A
^Hnginx-m2^RL^R

^HHäã<93><9a>^F^P^@^Z^@" ^A(^@2^A^ docker.io/library/nginx:latest:_docker.io/library/nginx@sha256:2f770d2fe27bc85f68fd7fe6a6390ef7076
BestEffortZ^@b^N
^L10.98.69.216^Z^@"^@

# 아까 생성한 nginx-m pod의 상세 정보 조회 > etcd가 저장한 pod의 정보와 유사한 부분들을 확인할 수 있음
$ kubectl describe po nginx-m2
Name: nginx-m2
Namespace: default
Status: Running
IP: 10.98.69.216
IPs:
IP: 10.98.69.216
Containers:
  nginx-m2:
    Container ID: containerd://6ff1306a4f878d831b6530f3a789ed0e4f00ec50962c698814fbffd12b841376
    Image: nginx
    Image ID: docker.io/library/nginx@sha256:2f770d2fe27bc85f68fd7fe6a6390ef7076bc703022fe81b980377fe3d27b70
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Tue, 11 Oct 2022 13:32:41 +0900
      Ready: True
```